

Anvesh Rao Vijjini

anveshrao1@gmail.com | ☎ 984-322-8184 |  anvesh-rao |  nvshrao | nvshrao.github.io

EDUCATION

- **University of North Carolina at Chapel Hill** 2021 to 2026
Advisor: Dr. Snigdha Chaturvedi Chapel Hill, NC
 - PhD in Computer Science
- **International Institute of Information Technology** 2013 to 2018
Advisor: Dr. Radhika Mamidi Hyderabad, India
 - MS in Computational Linguistics and B.Tech in Computer Science

EXPERIENCE

- **Microsoft** May 2025 to August 2025
Research Intern - Office of Applied Research (OAR), Mentors: Sheshera Mysore and Bahareh Sarrafzadeh
 - Engineered persona-driven multi-agent simulations and introduced novel metrics to quantify shifts in collaboration behaviour, such as disagreement rates following the deployment of LLM team assistants. Work [S.1] under review.
- **Bloomberg** May 2024 to July 2024
Research Intern, Mentors: Rajarshi Bhowmik and Jing Wang
 - Explored how LLMs encode rich entity representations in-context. Without any entity-specific fine-tuning, entity representations extracted from decoder-based LLMs can be applied to entity linking, entity typing, slot filling, and question answering by leveraging ranking and clustering.
- **Bloomberg** May 2023 to July 2023
Research Intern, Mentors: Junjie Liang and Shubham Kumar
 - Improved sparse document retrieval by 0.1 nDCG over the Doc2Query baseline by document expansion with LLMs-based question generation, with selection controlled by maximal marginal relevance (MMR).
- **Adobe** May 2022 to July 2022
Research Intern - Project Blink, Mentors: Hanieh Deilamsalehy and Franck Dernoncourt
 - Developed a pre-training strategy to improve transcript segmentation by pre-training in a curated order, with samples arranged in increasing similarity to the downstream task. Work published at EACL 2023 (Findings) and patented as U.S. Patent No. 12367344B2.
- **Samsung R&D** Aug 2018 to Aug 2021
Lead Engineer (Research) - Bixby Intelligence, (Samsung's voice assistant)
 - Automated WER Prediction of Bixby's ASR, saving 100s of hours of manual effort to listen to and transcribe user utterances for WER calculation.
 - Developed an automated Conversation Quality Estimation model to evaluate Bixby. This solution scaled evaluation capacity from hundreds of manual reviews to over 20K conversations per month.
 - Authored research papers based on Text classification and Domain Adaptation at EMNLP 2019, CICLing 2019, NLDB 2019, NLDB 2020, ICPR 2020 and EACL 2021.

PUBLICATIONS

C=CONFERENCE, W=WORKSHOP, S=UNDER REVIEW, P=PATENT, T=THESIS

- [S.1] **Evaluating Context Sensitivity of Collaboration Behavior in Multi-Party Simulation Frameworks.** Anvesh Rao Vijjini, Sheshera Mysore, Pegah Jandaghi, Ashish Sharma, Snigdha Chaturvedi, Bahareh Sarrafzadeh May 2026.
- [S.2] **Balancing Stereotype Exposure in Preference Data to Mitigate Personalization Bias.** Anvesh Rao Vijjini and Snigdha Chaturvedi. Jan 2026.
- [S.3] **PersonaMem-v2: Towards Personalized Intelligence via Learning Implicit User Personas and Agentic Memory.** Bowen Jiang, Yuan Yuan, Maohao Shen, Zhuoqun Hao, Zhangchen Xu, Zichen Chen, Ziyi Liu, Anvesh Rao Vijjini, Jiashu He, Hanchao Yu, Radha Poovendran, Gregory Wornell, Lyle Ungar, Dan Roth, Sihao Chen, Camillo Jose Taylor. ArXiv. Dec 2025.
- [C.1] **Do LLM Agents Mirror Socio-Cognitive Effects in Power-Asymmetric Conversations?** Anvesh Rao Vijjini and Snigdha Chaturvedi. ACL 2026.

- [C.2] **Exploring Safety-Utility Trade-Offs in Personalized Language Models.** Anvesh Rao Vijjini*, Somnath Basu Roy Chowdhury, and Snigdha Chaturvedi. NAACL 2025 (oral).
- [C.3] **SocialGaze: Improving the Integration of Human Social Norms in Large Language Models.** Anvesh Rao Vijjini, Rakesh Menon, Jiyai Fu, Shashank Shrivastava and Snigdha Chaturvedi. EMNLP 2024 (Findings).
- [C.4] **PARROT: Zero-Shot Narrative Reading Comprehension via Parallel Reading.** Chao Zhao, Anvesh Rao Vijjini, and Snigdha Chaturvedi. EMNLP 2023 (Findings).
- [C.5] **Curricular Next Conversation Prediction Pretraining for Transcript Segmentation.** Anvesh Rao Vijjini, Hanieh Deilamsalehy, Franck Dernoncourt and Snigdha Chaturvedi. EACL 2023 (Findings).
- [C.6] **Towards Inter-character Relationship-driven Story Generation.** Anvesh Rao Vijjini, Faeze Brahman and Snigdha Chaturvedi. EMNLP 2022.
- [C.7] **WER-BERT: Automatic WER Estimation with BERT in a Balanced Ordinal Classification Paradigm.** Anvesh Rao Vijjini*, Akshay Krishna Sheshadri* and Sukhdeep Kharbanda. EACL 2021.
- [C.8] **Sequential Domain Adaptation through Elastic Weight Consolidation for Sentiment Analysis.** Avinash Madasu and Anvesh Rao Vijjini. ICPR 2020.
- [C.9] **A SentiWordNet Strategy for Curriculum Learning in Sentiment Analysis.** Anvesh Rao Vijjini*, Kaveri Anuranjana* and Radhika Mamidi. NLDB 2020.
- [C.10] **Sequential Learning of Convolutional Features for Effective Text Classification.** Avinash Madasu and Anvesh Rao Vijjini. EMNLP-IJCNLP 2019.
- [C.11] **Gated Convolutional Neural Networks for Domain Adaptation.** Avinash Madasu and Anvesh Rao Vijjini. NLDB 2019.
- [C.12] **HindiRC: A Dataset for Reading Comprehension in Hindi.** Anvesh Rao Vijjini*, Kaveri Anuranjana* and Radhika Mamidi. CICLing 2019.
- [C.13] **Effectiveness of Self Normalizing Neural Networks for Text Classification.** Avinash Madasu and Anvesh Rao Vijjini. CICLing 2019.
- [C.14] **Twitter corpus of Resource-Scarce Languages for Sentiment Analysis and Multilingual Emoji Prediction.** Nurendra Choudhary, Rajat Singh, Anvesh Rao Vijjini and Manish Shrivastava. COLING 2018.
- [W.1] **Analyzing Curriculum Learning for Sentiment Analysis along Task Difficulty, Pacing and Visualization Axes.** Anvesh Rao Vijjini*, Kaveri Anuranjana* and Radhika Mamidi. Workshop on Computational Approaches to Subjectivity, Sentiment & Social Media Analysis (WASSA) at EACL 2021.
- [W.2] **Hindi Question Generation Using Dependency Structures.** Anvesh Rao Vijjini*, Kaveri Anuranjana* and Radhika Mamidi. Workshop on Humanizing AI (HAI) at IJCAI 2019.
- [W.3] **Towards Enhancing Lexical Resource and Using Sense-annotations of OntoSenseNet for Sentiment Analysis.** Sreekavitha Parupalli, Anvesh Rao Vijjini and Radhika Mamidi. Workshop on Semantic Deep Learning (SemDeep-3) at COLING 2018.
- [P.1] **Curricular Next Conversation Prediction Pretraining for Transcript Segmentation.** Inventors: Anvesh Rao Vijjini, Hanieh Deilamsalehy, Franck Dernoncourt. U.S. Patent No. 12367344B2 .
- [T.1] **Towards Adapting Curriculum Learning for Sentiment Analysis: Challenges and Analyses.** Anvesh Rao Vijjini and Radhika Mamidi. MS Thesis.

VOLUNTEER AND AWARDS

• Awards

Awards and Achievements

- Awarded 2026 Summer Research Fellowship at UNC (not accepted)
- Awarded Dean's List for being in the top 10% of the class in Fall 2015 and Spring 2016 at IIIT-H

• Volunteer & Service

Program committees, reviewing, and roles

- Reviewer: ACL 2020; EMNLP 2020; NLP+CSS @ EMNLP 2020; EACL 2021; ACL 2021; EMNLP 2022; NAACL 2023; EACL 2023; ACL 2023; WNU @ EMNLP 2024; ARR 2025 (Feb, May, Jul, Oct); ARR 2026 (Jan, Mar)
- Student Volunteer: EACL 2023
- Sessions Chair: NLDB 2020
- Social Media Chair: CoNLL 2025
- Board member: STEM Pride of the Triangle (stempride.web.unc.edu)